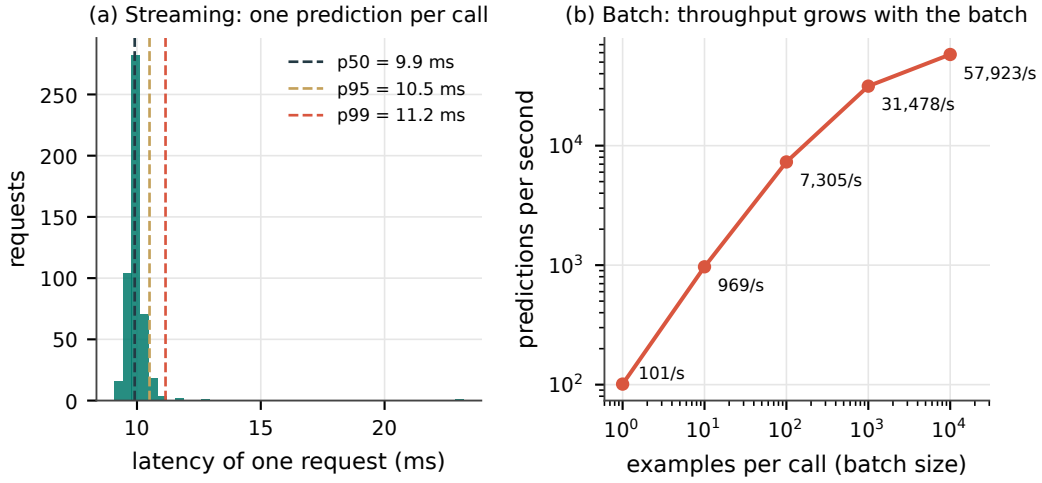
(a) Streaming: one prediction per call
p50 = 9.9 ms
p95 = 10.5 ms
p99 = 11.2 ms
requests
0
50
100
150
200
250
10
15
20
latency of one request (ms)
(b) Batch: throughput grows with the batch
predictions per second
10^2
10^3
10^4
101/s
969/s
7,305/s
31,478/s
57,923/s
10^0
10^1
10^2
10^3
10^4
examples per call (batch size)